

(Still a work-in progress; I want to revisit with intuitive explanations and playing-card examples)

Sorting is a key to CS theory, but easy to forget. I had an itch to review the algorithms in Wikipedia (strange, I know), and here are my notes:

## High-level thoughts

- Some algorithms (selection, bubble, heapsort) work by moving elements to their final position, one at a time. You sort an array of size  $N$ , put 1 item in place, and continue sorting an array of size  $N - 1$  (heapsort is slightly different).
- Some algorithms (insertion, quicksort, counting, radix) put items into a temporary position, close(r) to their final position. You rescan, moving items closer to the final position with each iteration.
- One technique is to start with a “sorted list” of one element, and merge unsorted items into it, one at a time.
- Complexity and running time
  - Factors: algorithmic complexity, startup costs, additional space requirements, use of recursion (function calls are expensive and eat stack space), worst-case behavior, assumptions about input data, caching, and behavior on already-sorted or nearly-sorted data
  - Worst-case behavior is important for real-time systems that need guaranteed performance. For security, you want the guarantee that data from an attacker does not have the ability to overwhelm your machine.
  - Caching — algorithms with sequential comparisons take advantage of spatial locality and prefetching, which is good for caching.
  - Algorithmic time vs. real time — The simple algorithms may be  $O(N^2)$ , but have low overhead. They can be faster for sorting small data sets ( $< 10$  items). One compromise is to use a different sorting method depending on the input size.
  - “Comparison sorts” make no assumptions on the data and compare all elements against each other (majority of sorts).  $O(N \lg N)$  time is the ideal “worst-case” scenario (if that makes sense —  $O(N \lg N)$  is the smallest penalty you can hope for in the worst case). Heapsort has this behavior.
  - $O(N)$  time is possible if we make assumptions about the data and don’t need to compare elements against each other (i.e., we know the data falls into a certain range or has some distribution).  $O(N)$  clearly is the minimum sorting time possible, since we must examine every element at least once (how can you sort an item you do not even examine?).

## Notes

- Assume we are sorting a list or array of  $N$  elements
- Once sorted, smaller items are on the left (first item) and larger items are on the right (last item)

## Bubble Sort [Best: $O(n)$ , Worst: $O(N^2)$ ]

Starting on the left, compare adjacent items and keep “bubbling” the larger one to the right (it’s in its final place). Bubble sort the remaining  $N - 1$  items.

- Though “simple” I found bubble sort nontrivial. In general, sorts where you iterate backwards (decreasing some index) were counter-intuitive for me. With bubble-sort, either you bubble items “forward” (left-to-right) and move the endpoint backwards (decreasing), or

bubble items “backward” (right-to-left) and increase the left endpoint. Either way, some index is decreasing.

- You also need to keep track of the next-to-last endpoint, so you don’t swap with a non-existent item.

## **Selection Sort [Best/Worst: $O(N^2)$ ]**

Scan all items and find the smallest. Swap it into position as the first item. Repeat the selection sort on the remaining  $N-1$  items.

- I found this the most intuitive and easiest to implement — you always iterate forward ( $i$  from 0 to  $N-1$ ), and swap with the smallest element (always  $i$ ).

## **Insertion Sort [Best: $O(N)$ , Worst: $O(N^2)$ ]**

Start with a sorted list of 1 element on the left, and  $N-1$  unsorted items on the right. Take the first unsorted item (element #2) and insert it into the sorted list, moving elements as necessary. We now have a sorted list of size 2, and  $N-2$  unsorted elements. Repeat for all elements.

- Like bubble sort, I found this counter-intuitive because you step “backwards”
- This is a little like bubble sort for moving items, except when you encounter an item smaller than you, you stop. If the data is reverse-sorted, each item must travel to the head of the list, and this becomes bubble-sort.
- There are various ways to move the item leftwards — you can do a swap on each iteration, or copy each item over its neighbor

## **Quicksort [Best: $O(N \lg N)$ , Avg: $O(N \lg N)$ , Worst: $O(N^2)$ ]**

There are many versions of Quicksort, which is one of the most popular sorting methods due to its speed ( $O(N \lg N)$  average, but  $O(N^2)$  worst case). Here’s a few:

Using external memory:

- Pick a “pivot” item
- Partition the other items by adding them to a “less than pivot” sublist, or “greater than pivot” sublist
- The pivot goes between the two lists
- Repeat the quicksort on the sublists, until you get to a sublist of size 1 (which is sorted).
- Combine the lists — the entire list will be sorted

Using in-place memory:

- Pick a pivot item and swap it with the last item. We want to partition the data as above, and need to get the pivot out of the way.
- Scan the items from left-to-right, and swap items greater than the pivot with the last item (and decrement the “last” counter). This puts the “heavy” items at the end of the list, a little like bubble sort.
- Even if the item previously at the end is greater than the pivot, it will get swapped again on the next iteration.
- Continue scanning the items until the “last item” counter overlaps the item you are examining – it means everything past the “last item” counter is greater than the pivot.
- Finally, switch the pivot into its proper place. We know the “last item” counter has an item

greater than the pivot, so we swap the pivot there.

- Phew! Now, run quicksort again on the left and right subset lists. We know the pivot is in its final place (all items to left are smaller; all items to right are larger) so we can ignore it.

Using in-place memory w/two pointers:

- Pick a pivot and swap it out of the way
- Going left-to-right, find an oddball item that is greater than the pivot
- Going right-to-left, find an oddball item that is less than the pivot
- Swap the items if found, and keep going until the pointers cross — re-insert the pivot
- Quicksort the left and right partitions
- Note: this algorithm gets confusing when you have to keep track of the pointers and where to swap in the pivot

Notes

- If a bad pivot is chosen, you can imagine that the “less” subset is always empty. That means we are only creating a subset of one item smaller each time, which gives us  $O(N^2)$  behavior in the worst case.
- If you choose the first item, it may be the smallest item in a sorted list and give worst-case behavior. You can choose a random item, or median-of-three (front, middle, end).
- Quicksort is *fast* because it uses spatial locality — it walks neighboring elements, comparing them to the pivot value (which can be stored in a register). It makes very effective use of caching.
- The pivot is often swapped to the front, so it is out of the way during the pivoting. Afterwards, it is swapped into place (with a pivot item that is less than or equal to it, so the pivot is preserved).
- The quicksort algorithm is complicated, and you have to pass left and right boundary variables

## Heapsort [Best/Avg/Worst: $O(N \lg N)$ ]

Add all items into a heap. Pop the largest item from the heap and insert it at the end (final position). Repeat for all items.

- Heapsort is just like selection sort, but with a better way to get the largest element. Instead of scanning all the items to find the max, it pulls it from a heap. Heaps have properties that allow heapsort to work in-place, without additional memory.
- Creating the heap is  $O(N \lg N)$ . Popping items is  $O(1)$ , and fixing the heap after the pop is  $\lg N$ . There are  $N$  pops, so there is another  $O(N \lg N)$  factor, which is  $O(N \lg N)$  overall.
- Heapsort has  $O(N \lg N)$  behavior, even in the worst case, making it good for real-time applications

## Counting sort [Best/Avg/Worst: $O(N)$ ]

Assuming the data are integers, in a range of 0-k. Create an array of size  $K$  to keep track of how many items appear (3 items with value 0, 4 items with value 1, etc). Given this count, you can tell the position of an item — all the 1's must come after the 0's, of which there are 3. Therefore, the 1's start at item #4. Thus, we can scan the items and insert them into their proper position.

- Creating the count array is  $O(N)$
- Inserting items into their proper position is  $O(N)$
- I oversimplified here — there is a summation of the counts, and a greatest-to-least ordering which keeps the sort stable.

## Radix sort [Best/Avg/Worst: $O(N)$ ]

Get a series of numbers, and sort them one digit at a time (moving all the 1000's ahead of the 2000's, etc.). Repeat the sorting on each set of digits.

- Radix sort uses counting sort for efficient  $O(N)$  sorting of the digits ( $k = 0 \dots 9$ )
- Actually, radix sort goes from least significant digit (1's digit) to most significant, for reasons I'll explain later (see CLRS book)
- Radix & counting sort are fast, but require structured data, external memory and do not have the caching benefits of quicksort.

## Actually doing the sorts

For practice, I wrote most of the sorts above in C, based on the psuedocode. Findings

- Even “easy” sorts like bubble sort get complicated with decrements, off-by-one errors,  $>$  vs  $\geq$  as you try to avoid walking off the end of the array with a swap.
- Mocking up the problem on paper is crucial, just like writing the code to swap items in a linked list. Don't have it all in your head.
- I found and fixed bugs in all of my initial sorts. Create a good test harness that makes it *easy* to test.
  - I separated my sorting routines into a DLL (I'm learning how to do Windows programming — it's pretty different from Unix)
  - I created a simple command-line .exe that took a list of numbers, turned them into an array, and called my sorting function, printing the result. This type of testing was encouraged by Kernighan — the tests are easy, do not require compilation (such as hard-coding a “testing” program)
- Because testing was easy, I made every test case I could think of: Pre-sorted forward, backwards, 1 element, 2 elements, even and odd items, etc.
- For debugging, I printed the intermediate array at each stage of the sort.